

Case study 8: AI tool vendor pitch

Case study 8: AI tool vendor pitch

SYNTHETIC Fictional vendor materials for training only.

Scenario

A vendor demo claims their humanitarian AI platform will:

- “Predict displacement 30 days ahead with 95% accuracy”
- “Auto-approve cash assistance for low-risk cases”
- “Train on your data to improve models for the sector”
- “Comply with all privacy laws globally”

Your organization has no AI procurement framework yet.

Your task (12 minutes)

1. List claims that need evidence.
2. Which claims raise red flags for humanitarian use?
3. What questions would you ask the vendor?
4. Recommend: adopt now, pilot with limits, or decline. Who decides?

Model answer

Evidence needed: Accuracy study methodology, false positive/negative rates by subgroup, legal compliance map (which laws?), data retention and training use, human override for cash decisions.

Red flags:

- Auto-approval for assistance (accountability, appeal rights)
- Sector-wide training on your data without consent
- Vague “all laws” compliance claim

Vendor questions:

- Where is data processed? Sub-processors?
- Can we audit model decisions affecting individuals?
- What is the incident response process?
- Who is liable for wrongful denial of aid?

Recommendation: Amber at best: limited pilot only after governance framework, DPIA, and no auto-approval of benefits. Red for immediate adoption.

Decision owner: Leadership with privacy, protection, and program input.

Common errors

- Buying on accuracy marketing without subgroup testing
- Piloting on live beneficiary data without safeguards
- No exit plan if vendor changes terms or pricing

Principle taught

Vendor due diligence: Humanitarian contexts require evidence, not slogans. Governance precedes procurement.